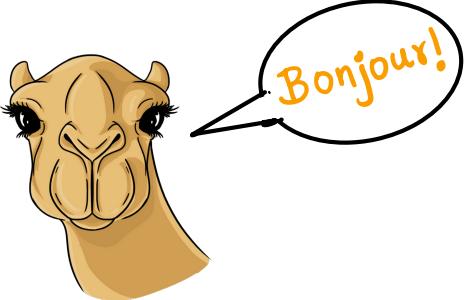


How Multicore GC Works



Meet our
old fren
Camel

He got cooler!
OCaml is finally
multicore \o/

OCaml came
into existance
in 1996
(older than me,
what?! :o)

OCaml stood the
test of time as a
stable and performant
FP language

Alas! OCaml
didn't support
multicore :(

For the longest
time, OCaml
didn't support
native parallelism

The challenge?
Building a multicore
capable **Garbage
Collector (GC)** that's
backwards compatible
and **performant**

In came a group
of hackers who
took up this
formidable task

They set on this
long and arduous
journey of building
a multicore GC

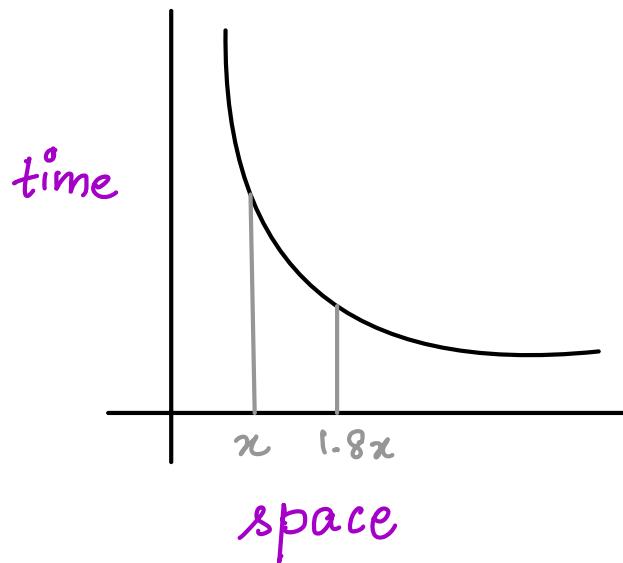
It finally got
merged and released
as OCaml 5.0 in
December 2022

Let's take a
deep dive on
the multicore
GC



Multicore logo

Space-time Trade off



- * operating a GC is essentially choosing a trade-off b/w Space and time
- * x is the theoretical min amount of space required
- * As the available space increases the program gets faster
- * Of course, one doesn't have the luxury of infinite space \Rightarrow $1.8x$ is usually * the sweet-spot *
- * depends on a multitude of factors

But, first things first - how are OCaml values represented in the memory?

MEMORY REPRESENTATION OF VALUES

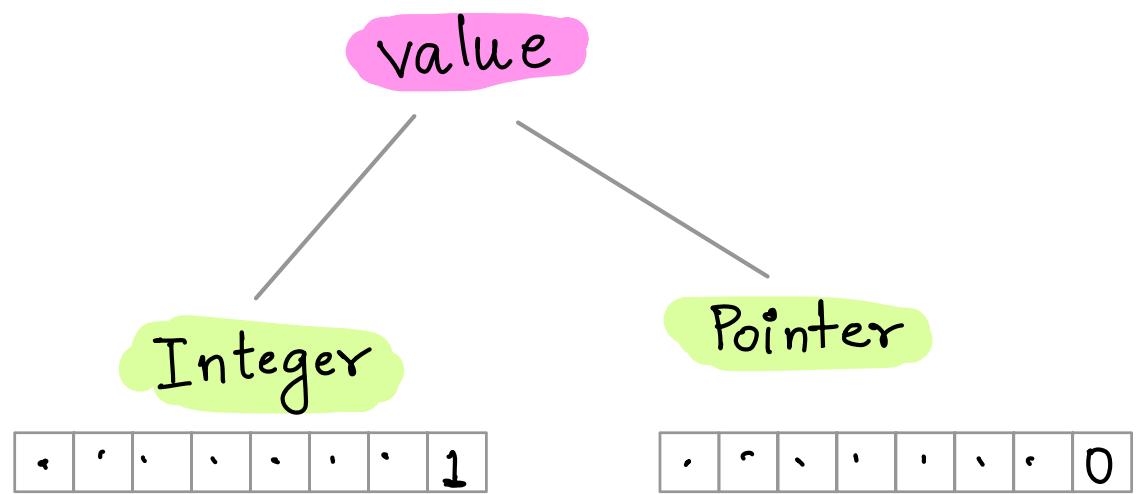


blocks of memory

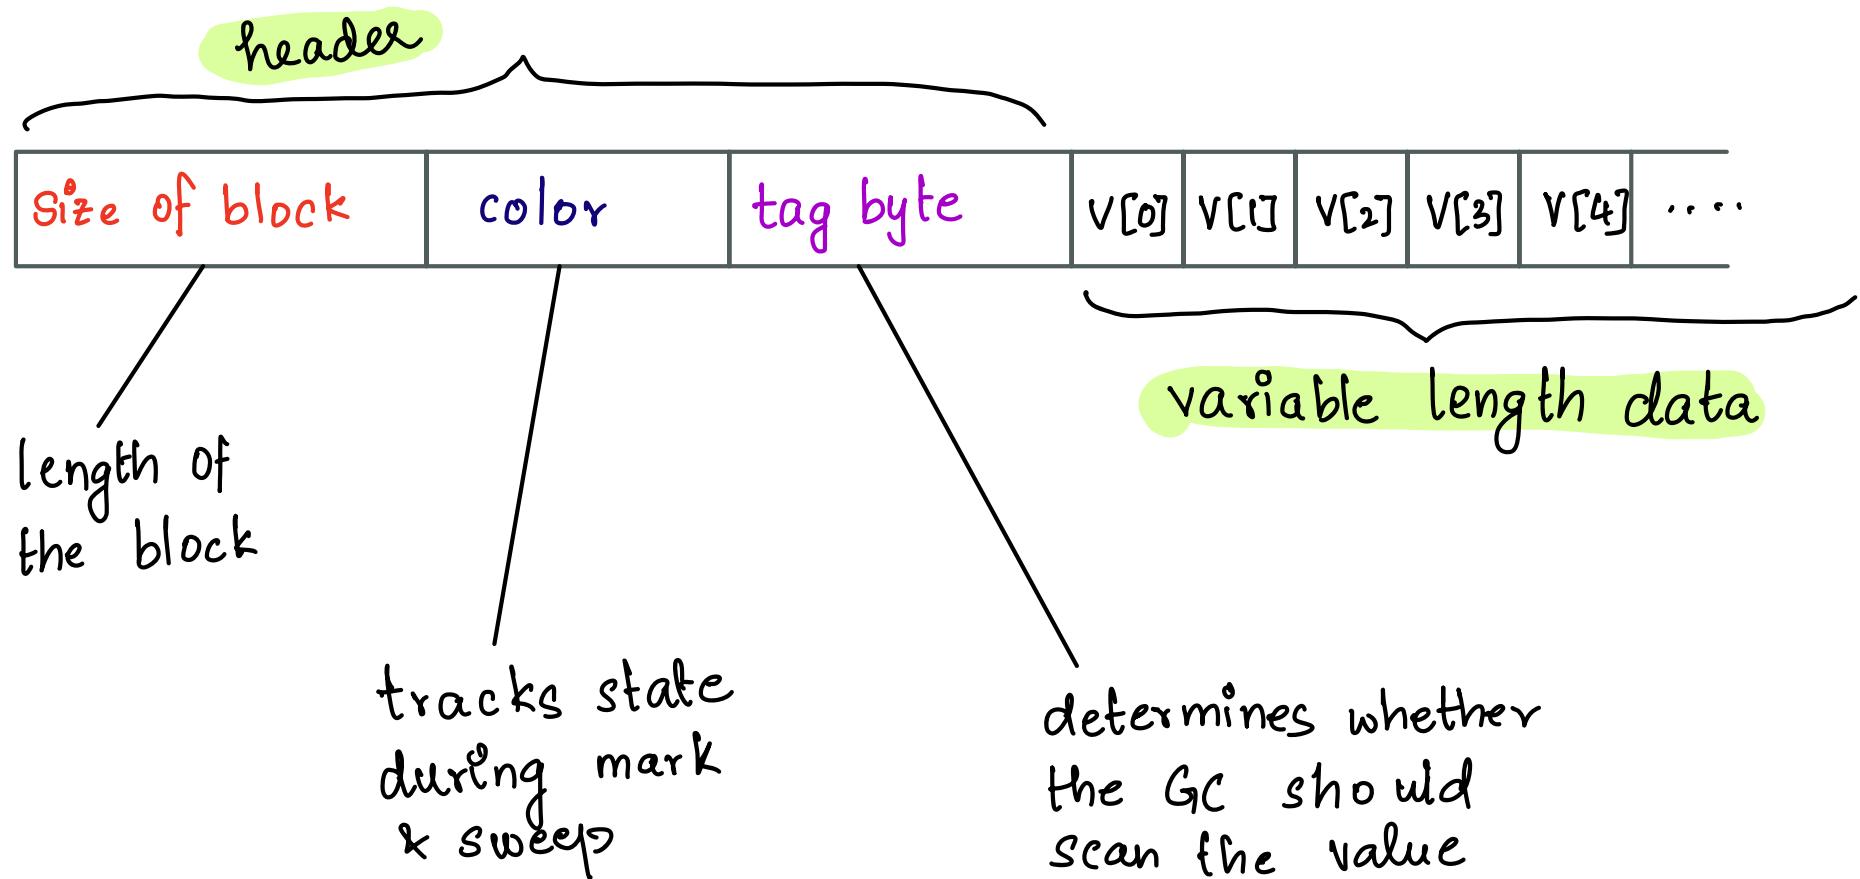
Values are either
integers (unboxed) or
pointers (boxed)

Boxed values have
meta-data attached
to it

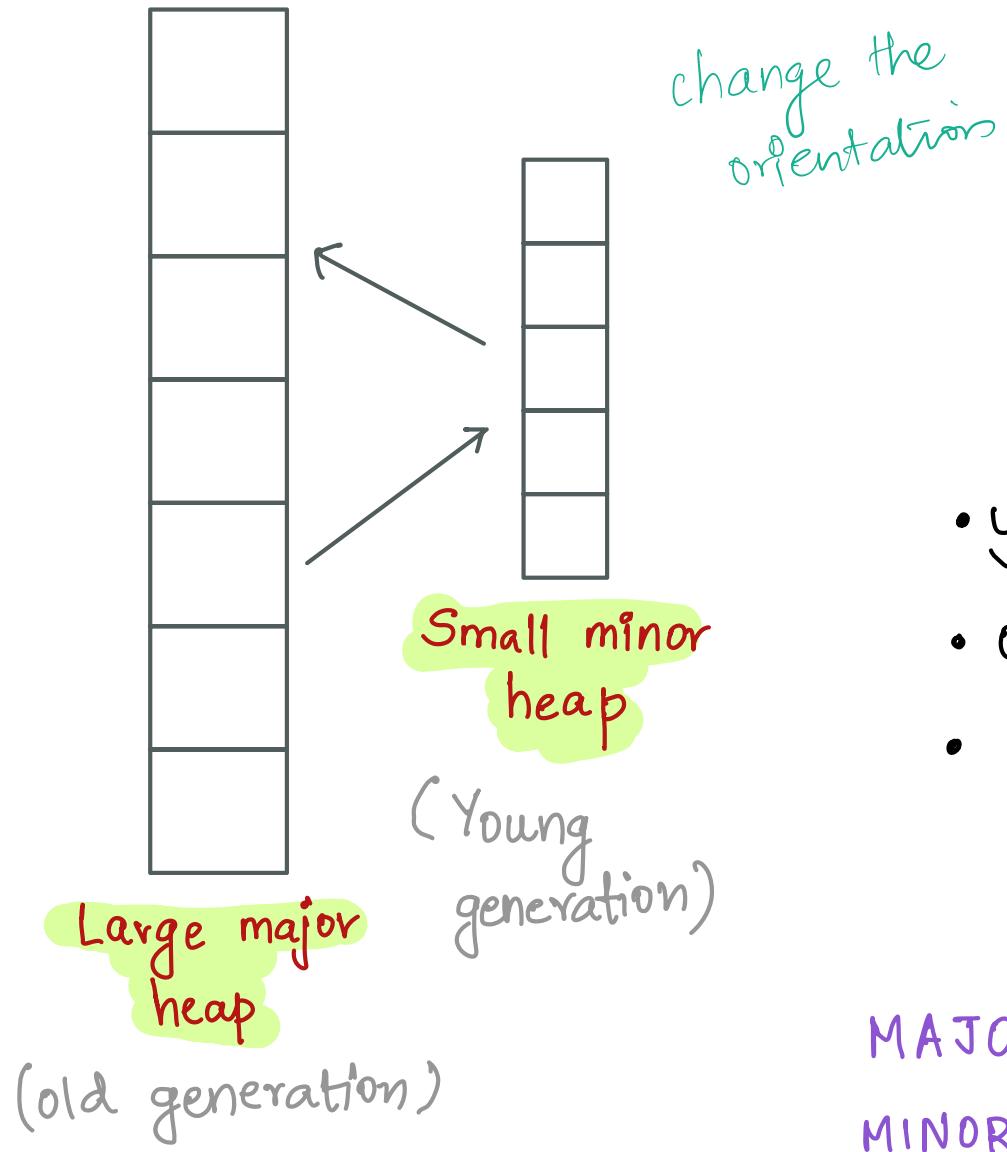
OCaml follows uniform memory layout ~ the layout is consistent regardless of type



Let's take a closer look at the constituents of a block:



GENERATIONAL HYPOTHESIS



- young blocks tend to die young
- old blocks tend to stay longer
- live blocks from minor heap are promoted to major heap

MAJOR HEAP: mark and sweep collection
MINOR HEAP: copying collection

KEY TAKEAWAYS

Allocations are fast: ~80% of allocations happen in minor heap

Pause times are small: Single digit ms pause times at max
GC memory is not exhausted

Survival rate is low: The survival rate of allocated objects
is roughly 10% on the minor heap.

Meaning, most allocated objects are
collected. Small objects die young

To speedup your program: Allocate as little as possible!



MINOR HEAP

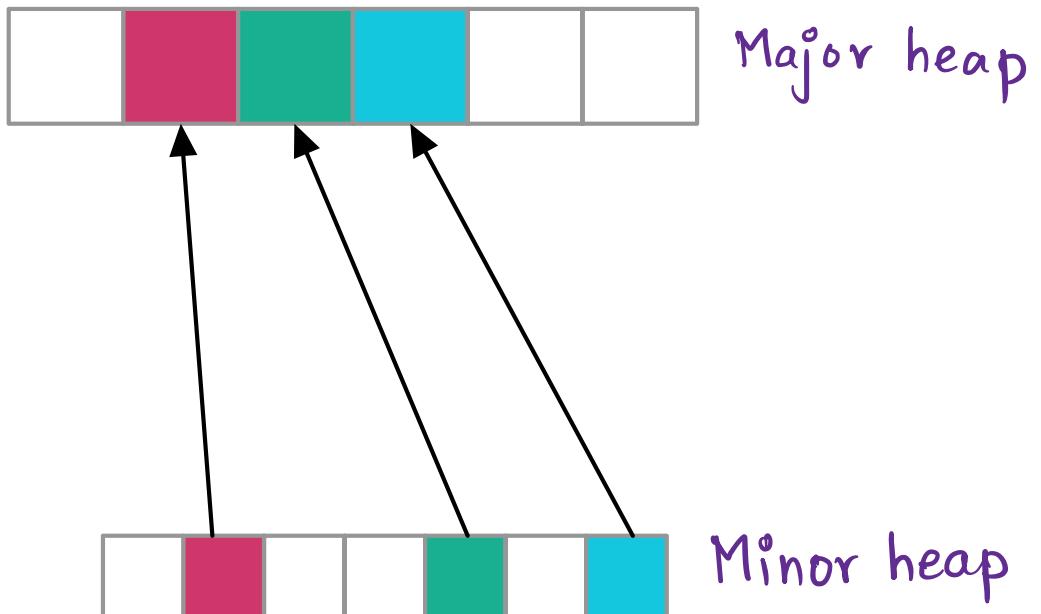
The minor heap is one single blob of memory



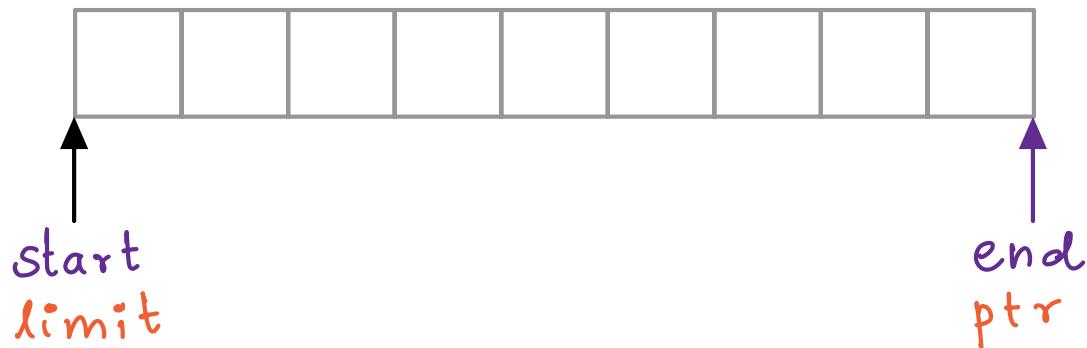
Live objects from minor heap are promoted to the major heap

Allocation is constant time and fast

Collection is non-incremental:
entire collection happens in one go



Allocation



Note: all the pointers have a 'caml_young' prefix. They're dropped for brevity

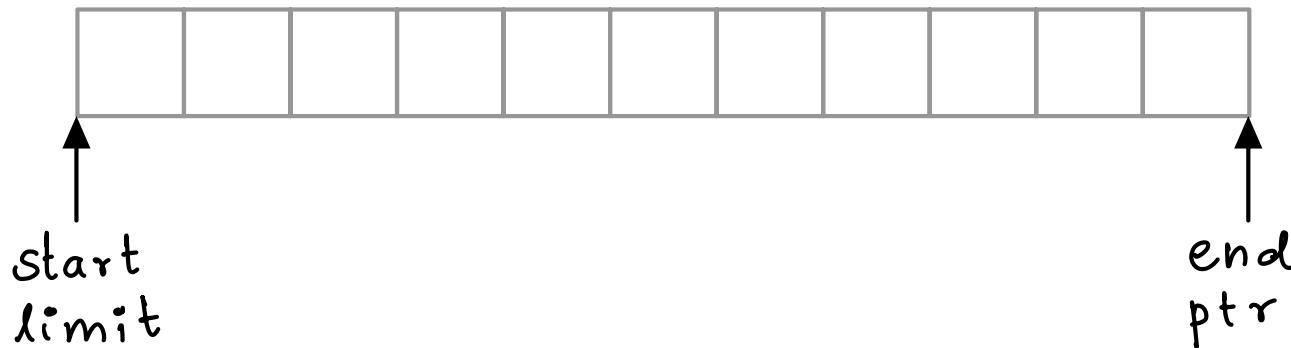
start: Starting address of the minor heap

end: finishing address of the minor heap

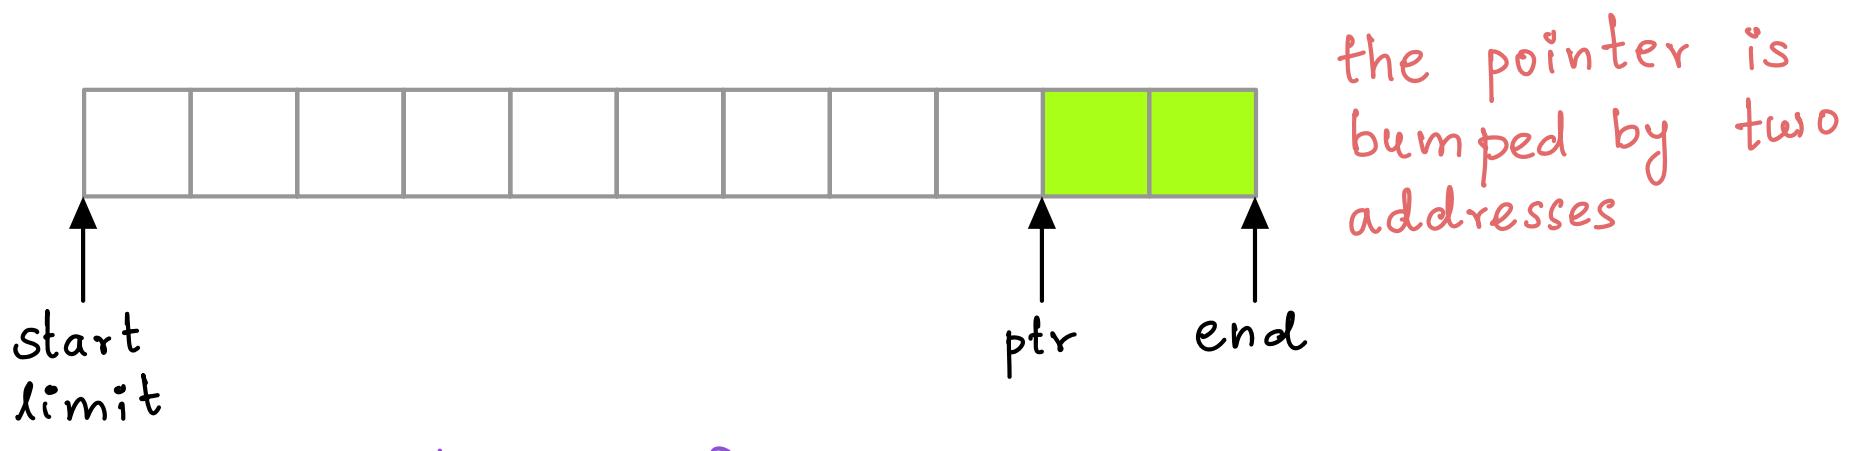
ptr: current location of the allocation pointer

limit: variable to define the end of minor heap.
making limit = end triggers a collection.

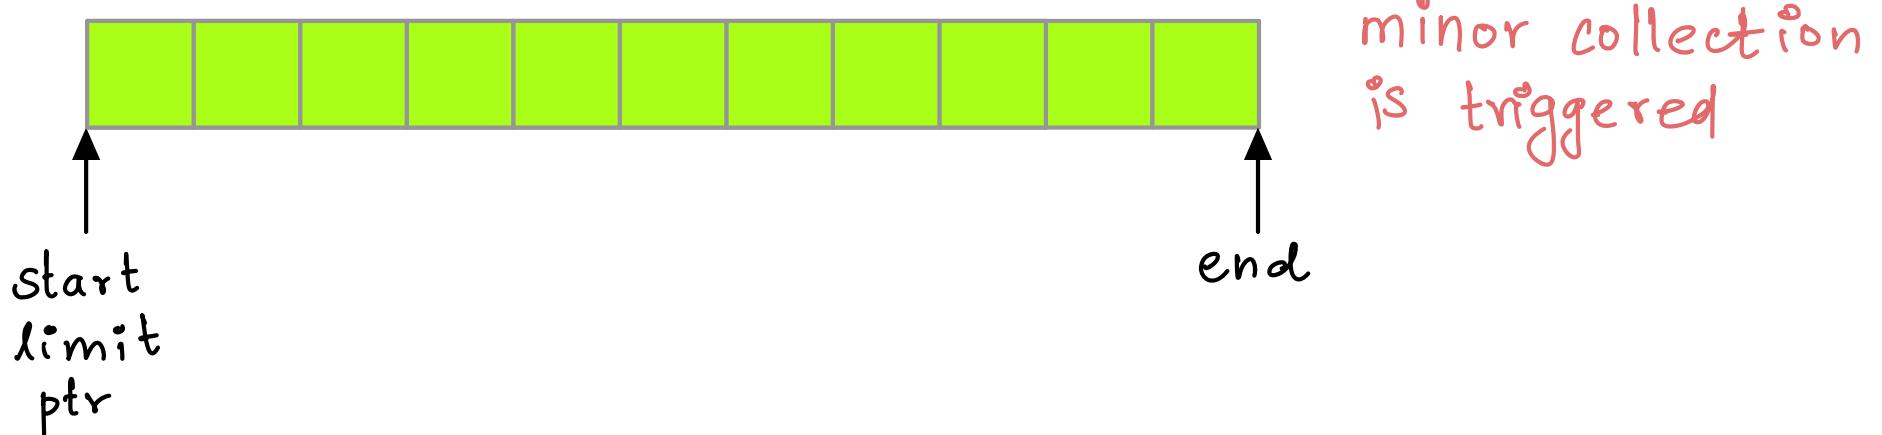
Initial state:



Allocate two blocks:



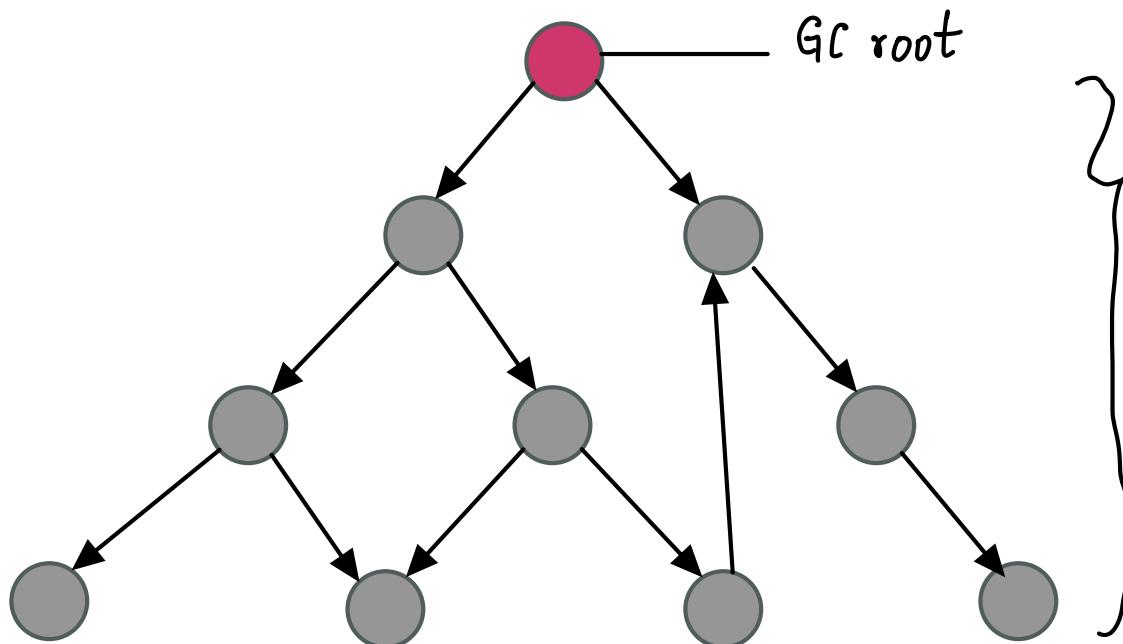
The minor heap is full:



MINOR COLLECTION

The collector traces live objects starting from a set of roots

Includes remembered set;
retains pointers from major heap to minor heap

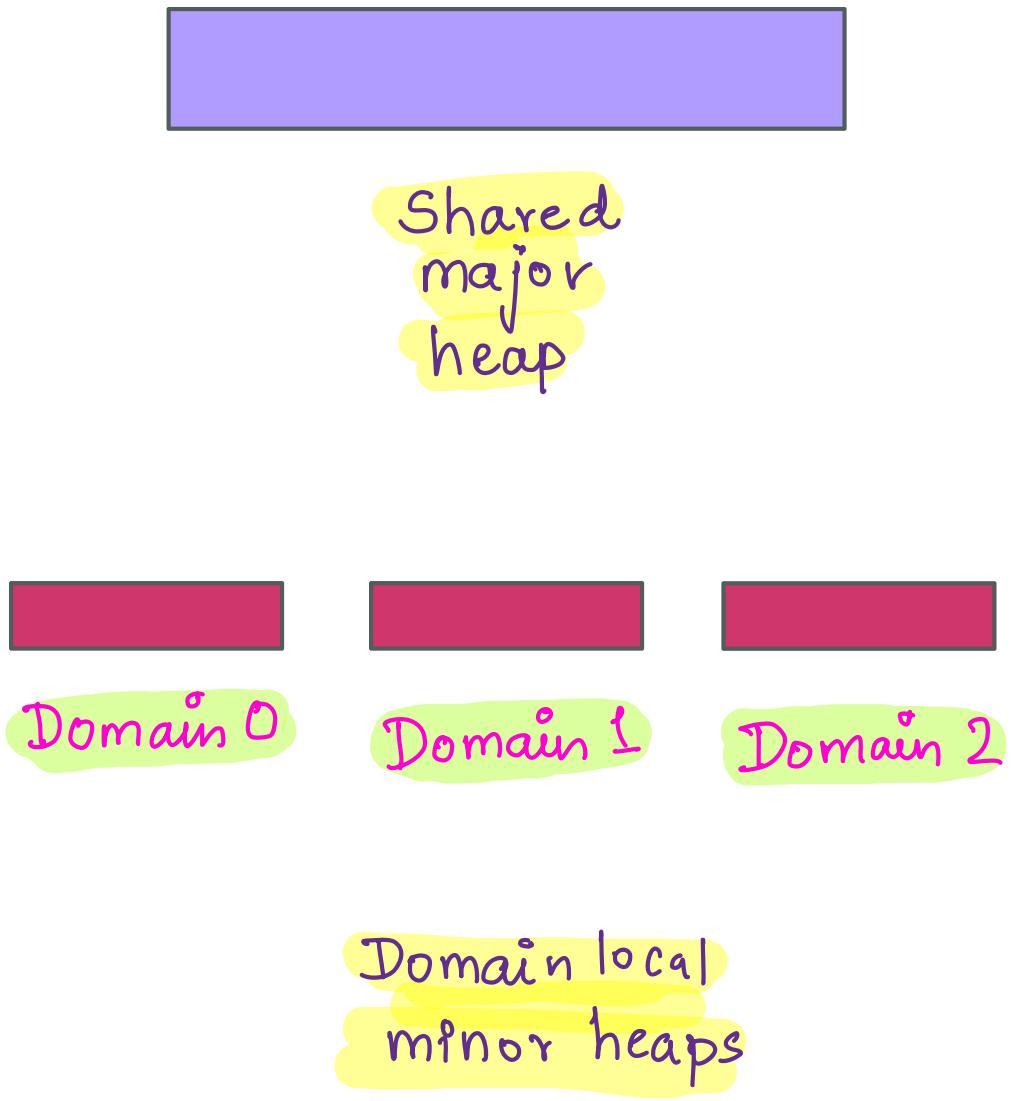


All reachable values from a root are traced via a **breadth-first search**

The live values are copied to the major heap (**promotion**)

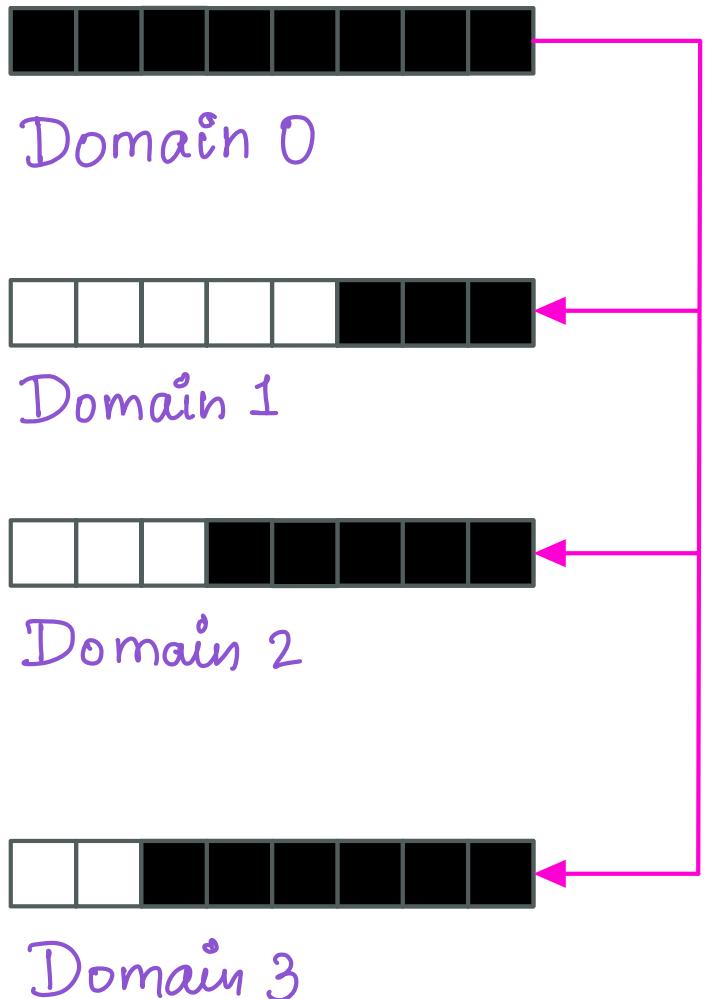
Garbage values are untouched and the **entire heap is cleared** at the end of a minor cycle

In comes multicore



- "Domain" is the basic unit of parallelism.
- The major heap is shared between domains.
- Every domain has its own minor heap
- A domain allocates in its own minor heap or the major heap
- Minor collection is stop-the-world and is performed by all the domains parallelly.

Stop-the-world Parallel Minor Collection



Domain 0
sends
interrupt
to other
domains
to trigger
minor
collection

Parallel promotion:
All domains promote live objects in parallel.
Two domains attempting to promote the same object is serialized.

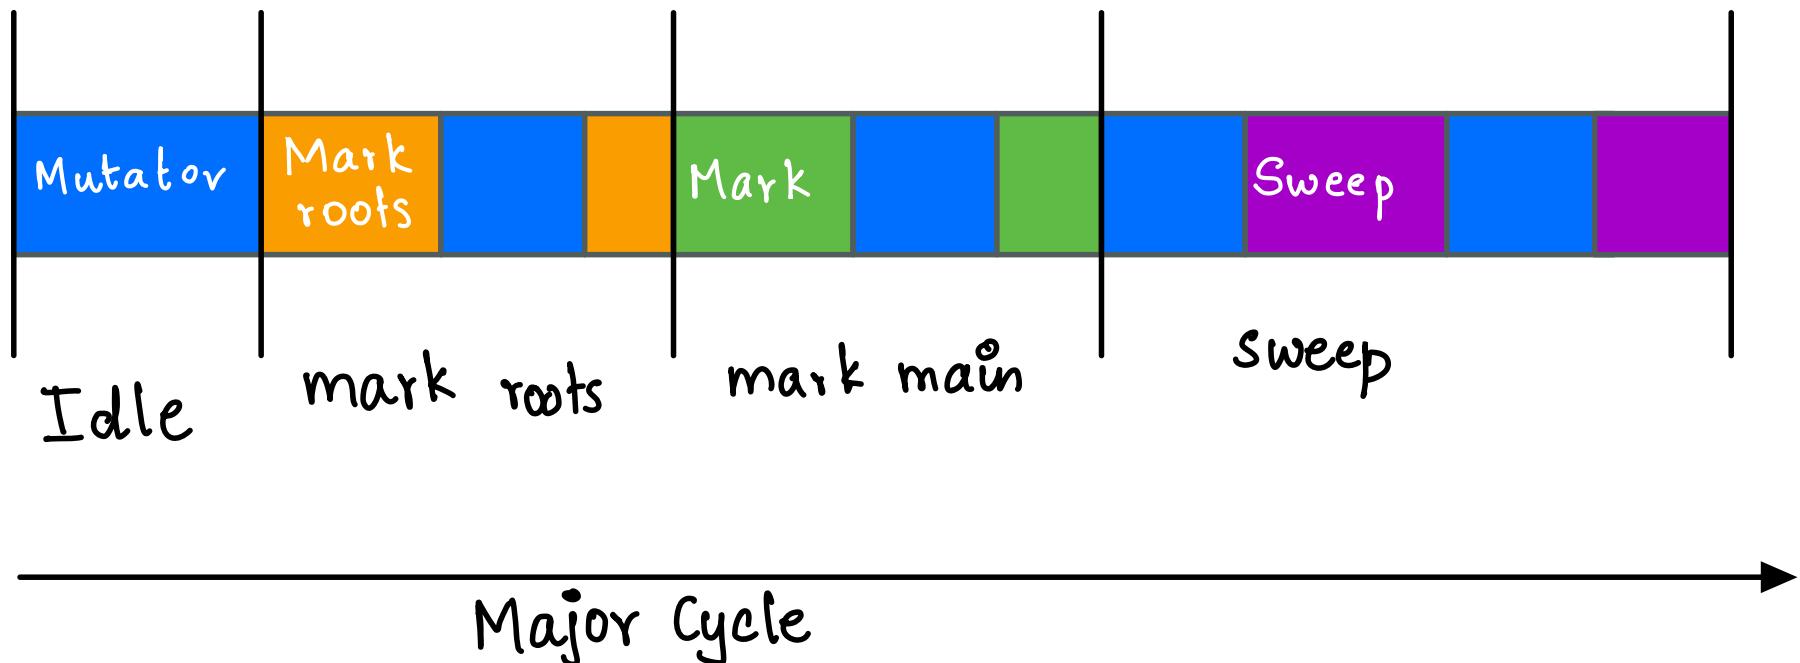
Static work-sharing:
The number of roots to scan is evenly divided amongst the domains. One domain could end up with more work.

MAJOR HEAP

Mark-and-Sweep: Marking determines allocations still in use.
Sweeping collects unused allocations.

Non-moving: Objects are in the same memory address.

Incremental: Program stops in small slices for low pause times



Marking

Consider a sequential program in OCaml 5

GC Colours:

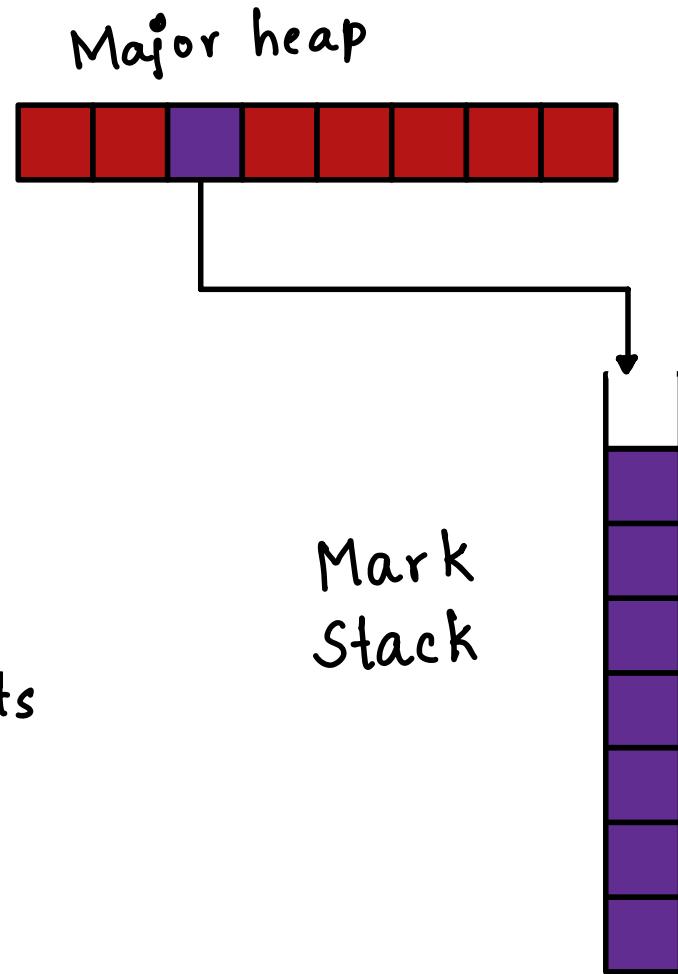
MARKED

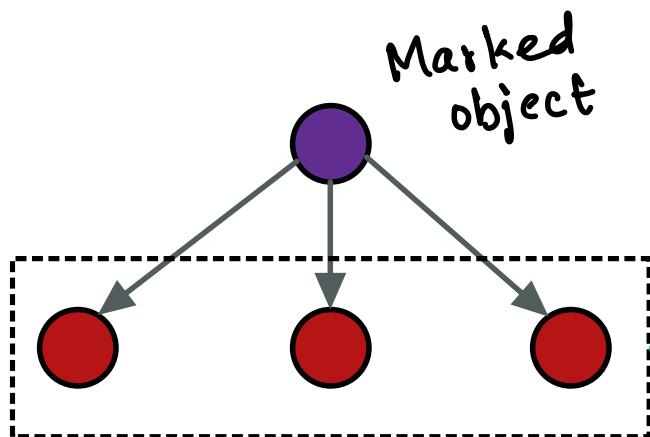
UNMARKED

GARBAGE

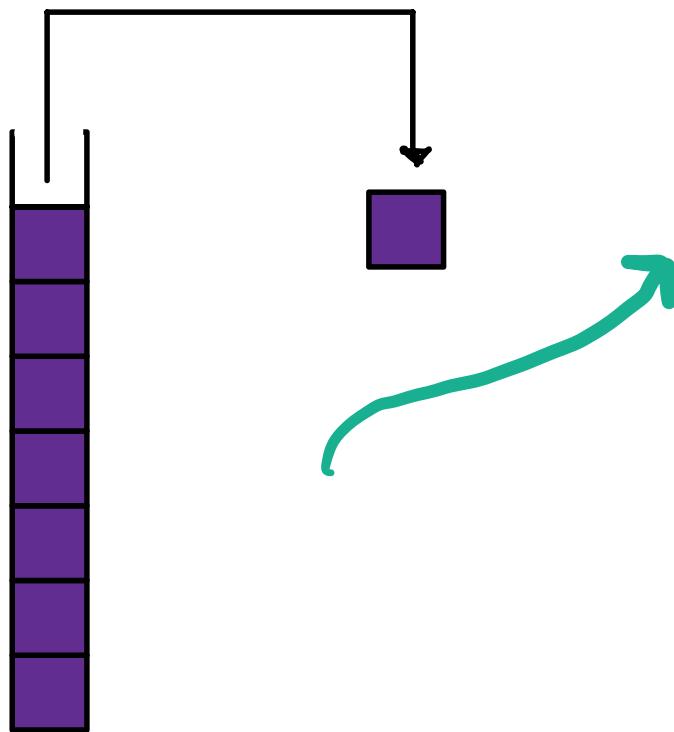
FREE

During the marking phase all live objects are marked and put into a mark stack





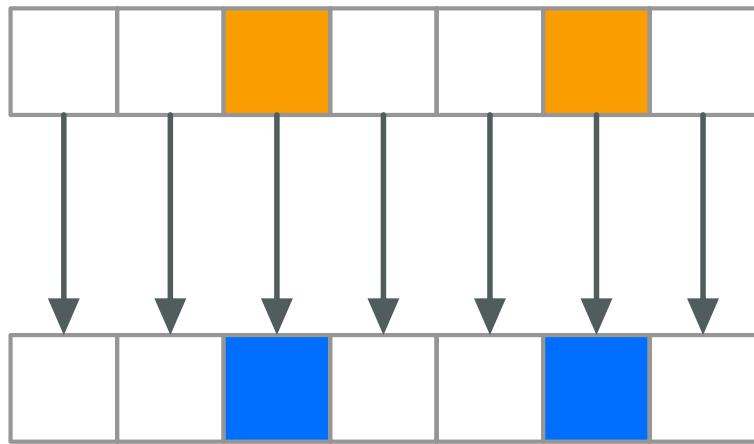
The objects **reachable** from a live object are also marked and pushed into a mark stack



Items from mark stack are popped and all its **children** are marked

This process is repeated until the mark stack is **empty**

Sweeping



Set of objects affected by marking and sweeping is disjoint. No synchronization is necessary.

Sweeping pass traverses the whole heap incrementally and flips the meaning of the state

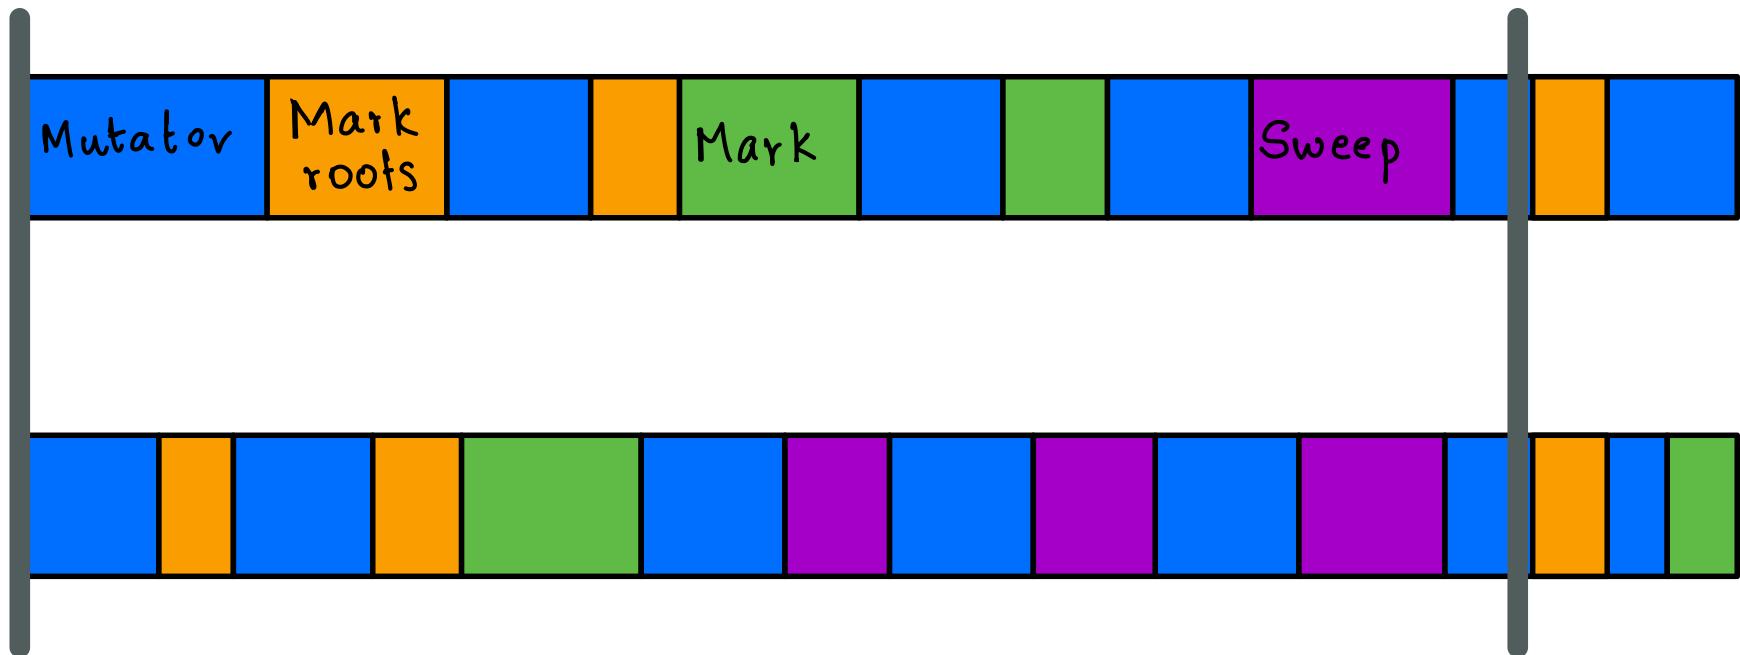
MARKING

Unmarked \Rightarrow Marked

SWEEPING

Garbage \Rightarrow Free

Multicore Major Heap



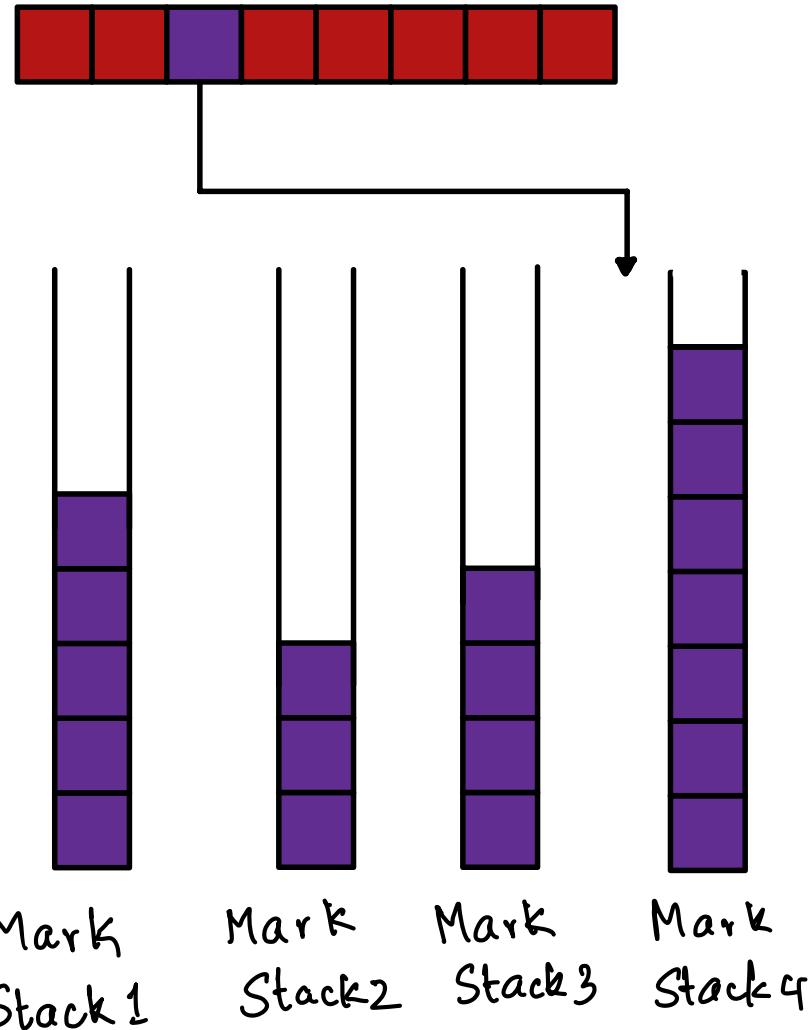
Start of
major cycle

The major collector's design allows overlap of mark & sweep phase, thus removing the need for synchronization.

End of major cycle

There's a small STW section at the end of a major cy

Major heap



- Every domain has its own mark stack
- Marking is idempotent: okay to have the same object on more than one mark stack.
- Sweeping is disjoint: Every domain only sweeps the memory it allocated.